

Sample estate: RECON_TEST on Oracle XE 21c

Schema **RECON_TEST** · Oracle 21.0.0.0.0 · extracted 2026-09-02 15:53:31
Prepared by DevCrafter · 2026-09-03 · engine: pgrecon



Executive summary

This assessment covers an Oracle 21.0.0.0.0 environment comprising 33 database objects and approximately 150 lines of PL/SQL. Analysis identified 58 discrete findings relevant to a PostgreSQL migration, spanning schema design, PL/SQL constructs, data types, and operational dependencies. The overall program has been graded C, indicating a major undertaking for which phased delivery is essential rather than optional.

Of the 58 findings, 10 are classified as high severity, 18 as medium, 16 as low, and 14 as informational. The high-severity items include a database link, interval partitioning, a global index on a partitioned table, package-level state used in two places, UTL_FILE usage, a compound trigger, an autonomous transaction, a LONG or LONG RAW column, and a BFILE column. These represent constructs with no direct PostgreSQL equivalent or that require deliberate redesign, and they should anchor the remediation plan. The medium and low findings, while individually less disruptive, are numerous and touch a broad range of areas including scheduler jobs, materialized views, dynamic SQL, collection types, LOB handling, and data type precision, all of which will require systematic review across the object inventory.

58 FINDINGS	10 HIGH	18 MEDIUM	131 EXPECTED PERSON-DAYS
-----------------------------------	-------------------------------	---------------------------------	--

Effort range: **106–180** person-days (5.1–8.6 person-months), expected 131 — including testing and stabilization. Assumptions are listed with the estimate; this is a scoping

instrument to be calibrated per engagement, not a quote.

Estate overview

33 OBJECTS	150 PL/SQL LINES	13/13 DDL PARSED	9 UNITS DEEP-PARSED
----------------------	----------------------------	----------------------------	-------------------------------

OBJECT TYPE	COUNT
TABLE	10
INDEX	8
VIEW	2
TRIGGER	2
PROCEDURE	2
TYPE BODY	1
TYPE	1
SYNONYM	1
SEQUENCE	1
PACKAGE BODY	1
PACKAGE	1
MATERIALIZED VIEW	1
JOB	1
FUNCTION	1

Largest tables by rows (from optimizer statistics)

TABLE	ROWS
MV_DEPT_SALARIES	3

License exposure

What the database reports about its own edition, hardware, and use of separately licensed features. These are the source system's facts; contract interpretation and pricing belong in the engagement,

next to the actual agreements.

Express

EDITION

4

CPU CORES

2

CPU_COUNT IN FORCE

No separately licensed option shows usage in the database's own statistics: the estate runs on the edition's base feature set.

Migration hot spots

The ten objects concentrating the most remediation weight; these decide the project's shape and deserve design attention first.

OBJECT	FINDINGS	WEIGHT	DRIVEN BY
RECON_TEST.PKG_LEDGER	6	12.3	Package-level state (5), Package initialization block (2), PL/SQL collection type (2), +3 more
RECON_TEST.ARCHIVE_NOTES	8	8.3	UTL_FILE usage (3), ROWID in code (1.5), DBMS_LOB usage (1.5), +5 more
RECON_TEST.SALES	4	6.8	Interval partitioning (6), Parallel degree set on object (0.5), DDL parsed only as an opaque command (0.3), +1 more
RECON_TEST.TRG_EMP_AUDIT	2	5.5	Autonomous transaction (4), Transaction control inside PL/SQL (1.5)
RECON_TEST.PKG_LEDGER (body)	1	5.0	Package-level state (5)
RECON_TEST.TRG_EMP_SAL_GUARD	2	4.3	Compound trigger (4), SYSDATE (0.3)
RECON_TEST.LEGACY_NOTES.BODY	1	4.0	LONG or LONG RAW column (4)
RECON_TEST.MV_DEPT_SALARIES	2	3.5	Materialized view (2), Query-rewrite materialized view (1.5)
RECON_TEST.LOOPBACK	1	3.0	Database link (3)
RECON_TEST.LEGACY_REFS.SCAN_DOC	1	3.0	BFILE column (3)

Findings and remedies

HIGH

Database link R-OBJ-01 • extension: postgres_fdw

Cross-database access needs postgres_fdw or oracle_fdw on the PostgreSQL side, plus credential management. Every statement using the link must be found and repointed.

RECON_TEST.LOOPBACK	to //localhost:1521/XEPDB1 as RECON_TEST
---------------------	--

HIGH

Interval partitioning R-PART-01 • extension: pg_partman

PostgreSQL has no partitioning scheme that creates partitions on data arrival. Plan pg_partman or scheduled partition creation, and decide the retention window explicitly.

RECON_TEST.SALES	INTERVAL NUMTOYMINTERVAL(1, 'MONTH')
------------------	--------------------------------------

HIGH

Global index on a partitioned table R-PERF-02

PostgreSQL indexes on partitioned tables are always local: there is no global index, and a unique key must contain the partition key. Global-index lookups become per-partition probes, and unique constraints that cannot include the key force a design decision, not a translation.

RECON_TEST.SALES_AMOUNT_GIX	GLOBAL partitioned index on SALES
-----------------------------	-----------------------------------

HIGH

Package-level state R-PKG-01

PostgreSQL has no package variables. Scalars map to session settings via set_config and current_setting (text only); composites and collections need per-session temporary tables. Every reader and writer of the state changes.

RECON_TEST.PKG_LEDGER	2 package-level declaration(s), first at line 2
RECON_TEST.PKG_LEDGER (body)	1 package-level declaration(s), first at line 2

HIGH **UTL_FILE usage** R-SYS-01 • extension: orafce

PostgreSQL functions have no supported server-side file API. Move file handling to the application, or accept the operational cost of an untrusted language such as plpython.

RECON_TEST.ARCHIVE_NOTES

UTL_FILE (first at line 16)

HIGH **Compound trigger** R-TRG-01

PostgreSQL has no compound triggers. Split the timing sections into separate trigger functions and move statement-level state into a transition table or a session variable.

RECON_TEST.TRG_EMP_SAL_GUARD

on EMP

HIGH **Autonomous transaction** R-TRG-02

PostgreSQL functions cannot commit independently of their caller. Rework audit-style autonomous transactions with a dblink loopback, a message table drained by a worker, or pg_background.

RECON_TEST.TRG_EMP_AUDIT

PRAGMA AUTONOMOUS_TRANSACTION (first at line 5)

HIGH **LONG or LONG RAW column** R-TYPE-01

LONG has been deprecated since Oracle 8. Map LONG to text and LONG RAW to bytea; any code reading the column piecewise must be rewritten, and only one LONG per table exists to find.

RECON_TEST.LEGACY_NOTES.BODY

LONG

HIGH **BFILE column** R-TYPE-07

BFILE points at a file outside the database and PostgreSQL has no counterpart. Store a path or object key plus the file in application storage, and move the DIRECTORY grants and BFILENAME plumbing into the application.

RECON_TEST.LEGACY_REFS.SCAN_DOC

BFILE

MEDIUM Scheduler or legacy job R-OBJ-02 • extension: pg_cron

DBMS_SCHEDULER and DBMS_JOB have no engine equivalent in PostgreSQL. Recreate schedules in pg_cron and port each job action; calendar syntax must be translated by hand.

RECON_TEST.NIGHTLY_ROLLUP

scheduler job

MEDIUM Materialized view R-OBJ-03

PostgreSQL materialized views refresh only in full and have no refresh-on-commit. Fast-refresh views need a trigger-based design or a scheduled REFRESH with the staleness accepted.

RECON_TEST.MV_DEPT_SALARIES

materialized view

MEDIUM User-defined object type R-OBJ-05

Composite types carry the attributes, but member functions do not exist in PostgreSQL: each method becomes a standalone function taking the composite as its first argument, and every call site changes.

RECON_TEST.T_MONEY

object type

MEDIUM Query-rewrite materialized view R-PERF-05

Oracle rewrites matching queries to hit the materialized view transparently; PostgreSQL never does. Queries that benefited must reference the view explicitly, which means finding them in the application first.

RECON_TEST.MV_DEPT_SALARIES

REWRITE enabled, refresh COMPLETE

MEDIUM Package initialization block R-PKG-02

Oracle runs the block once per session on first reference, which PostgreSQL cannot replicate. Extract it into an init function guarded by a session flag and call it lazily from each entry point. Detection is heuristic; confirm by eye.

RECON_TEST.PKG_LEDGER

BEGIN blocks exceed subprograms by 1 (possible initialization block)

MEDIUM Dynamic SQL R-SRC-01

EXECUTE IMMEDIATE maps to PL/pgSQL EXECUTE, but bind syntax and error handling differ. Each dynamic statement needs a manual port and a test.

RECON_TEST.PKG_LEDGER

EXECUTE IMMEDIATE (first at line 27)

MEDIUM GOTO statement R-SRC-02

PL/pgSQL has no GOTO. The control flow must be restructured with loops, exceptions, or early returns.

RECON_TEST.FIND_MANAGER

GOTO (first at line 14)

MEDIUM BULK COLLECT R-SRC-03

PL/pgSQL has no BULK COLLECT; array_agg into an array or a plain set-based statement usually replaces the whole loop.

RECON_TEST.PKG_LEDGER

BULK COLLECT (first at line 24)

MEDIUM CONNECT BY hierarchical query R-SRC-04

Rewrite as a recursive CTE (WITH RECURSIVE). LEVEL, SYS_CONNECT_BY_PATH, and ORDER SIBLINGS BY all need manual equivalents.

RECON_TEST.V_ORG_CHART

CONNECT BY in view definition

MEDIUM Transaction control inside PL/SQL R-SRC-14

PostgreSQL functions cannot COMMIT or ROLLBACK; only procedures called outside a transaction block can. Each routine that controls transactions must become a procedure or hand that control back to the caller.

RECON_TEST.TRG_EMP_AUDIT

COMMIT or ROLLBACK (first at line 15)

MEDIUM PL/SQL collection type R-SRC-16

Nested tables, varrays, and associative arrays map to PostgreSQL arrays, jsonb, or temporary tables depending on use; collection methods such as EXTEND, COUNT, and FIRST have no direct equivalents and every reader changes.

RECON_TEST.PKG_LEDGER

collection type declaration (first at line 20)

MEDIUM Empty string treated as NULL R-SRC-18

Oracle treats "" as NULL; PostgreSQL treats it as a real value. `v := ''` assigns NULL in Oracle but an empty string in PostgreSQL, `NVL(x, '')` stops collapsing, and `col = ''` can start matching rows. Nothing errors: behavior changes silently. Decide per site whether NULL or empty was meant, and put the data path under test.

RECON_TEST.ARCHIVE_NOTES

empty-string literal (first at line 4)

MEDIUM ROWID in code R-SRC-19

ctid is not a ROWID: it moves on update and vacuum and must not be stored or compared. Fetch-by-rowid loops rewrite to primary-key access or WHERE CURRENT OF; ROWID-typed variables and parameters change type with them.

RECON_TEST.ARCHIVE_NOTES

ROWID (first at line 3)

MEDIUM DBMS_LOB usage R-SYS-04

Piecewise LOB calls mostly disappear: text and bytea are handled whole, with substring and length built in. Streaming access beyond 1 GB needs large objects and the `lo_` API.

RECON_TEST.ARCHIVE_NOTES

DBMS_LOB (first at line 14)

MEDIUM Global temporary table R-TAB-01 • extension: pgtt

PostgreSQL temporary tables are per session and vanish with it; there is no shared definition with private rows. Rework to on-the-fly temp tables or use the pgtt extension.

RECON_TEST.STAGING_ROWS

GLOBAL TEMPORARY

MEDIUM XMLTYPE column R-TYPE-02

PostgreSQL xml stores documents but has no XML index type and none of the XMLType member functions; XQuery support differs. Inventory how the application queries the XML before mapping.

RECON_TEST.PRODUCT_SPECS.SPEC

XMLTYPE

MEDIUM TIMESTAMP WITH LOCAL TIME ZONE column R-TYPE-03

Values are normalized to the session time zone on read, which timestamptz does not replicate exactly. Verify every reader assumes UTC or one fixed zone before mapping to timestamptz.

RECON_TEST.INVOICES.CREATED_AT

TIMESTAMP(6) WITH LOCAL TIME ZONE

MEDIUM ROWID or UROWID column R-TYPE-06

A stored ROWID is a physical row address, and PostgreSQL has no stable one: ctid changes on update and vacuum. Replace the column with a real key reference; whatever wrote and followed these addresses changes with it.

RECON_TEST.LEGACY_REFS.REF_ADDR

ROWID

LOW Object grants to migrate R-ENV-02

The converter emits role bootstraps and GRANT statements for privileges with a direct counterpart; EXECUTE grants and Oracle-only privileges are residue lines. Review the role list against the real user model before applying.

RECON_TEST.PUBLIC

1 object grants

LOW Synonym R-OBJ-06

PostgreSQL has no synonyms. Same-database synonyms usually dissolve into search_path; cross-schema renames become views or schema qualification at the call sites.

RECON_TEST.STAFF

for RECON_TEST.EMP

LOW **Optimizer hint** R-PERF-01 • extension: pg_hint_plan

PostgreSQL ignores optimizer hints; pg_hint_plan can pin plans but is a last resort. A hinted statement was tuned by hand against Oracle's planner, so budget a fresh look at its plan and indexes on PostgreSQL instead of a plain rewrite.

RECON_TEST.ARCHIVE_NOTES

optimizer hint (first at line 8)

LOW **Parallel degree set on object** R-PERF-03

PostgreSQL has no per-object parallel degree; the planner decides within max_parallel_workers_per_gather. Workloads that relied on forced parallel scans need a plan review and possibly session-level parallel settings.

RECON_TEST.SALES

DEGREE 4

LOW **Old-style outer join (+)** R-SRC-05

The (+) operator does not exist in PostgreSQL. Rewrite as ANSI LEFT or RIGHT JOIN; semantics match except for a few multi-condition corner cases worth testing.

RECON_TEST.V_EMP_DEPT

uses (+) outer join

LOW **SYSDATE** R-SRC-11

SYSDATE is evaluated at call time; now() is fixed at transaction start. Map to statement_timestamp() or clock_timestamp() where the difference matters, for example inside long transactions and loops.

RECON_TEST.ARCHIVE_NOTES

SYSDATE (first at line 19)

RECON_TEST.PKG_LEDGER

SYSDATE (first at line 7)

RECON_TEST.TRG_EMP_SAL_GUARD

SYSDATE (first at line 8)

LOW **DBMS_OUTPUT usage** R-SYS-05 • extension: orafce

Map PUT_LINE to RAISE NOTICE for diagnostics. If callers consume the output programmatically, that pattern needs a real interface instead.

RECON_TEST.ARCHIVE_NOTES

DBMS_OUTPUT (first at line 19)

LOW NUMBER without precision R-TYPE-04

Bare NUMBER accepts up to 38 significant digits. Map to numeric, and check whether application code assumes integer behavior that a typed bigint would express better.

RECON_TEST.AUDIT_LOG.AUDIT_ID	NUMBER with no precision
RECON_TEST.INVOICES.INVOICE_ID	NUMBER with no precision
RECON_TEST.MV_DEPT_SALARIES.HEADCOUNT	NUMBER with no precision
RECON_TEST.MV_DEPT_SALARIES.TOTAL_SAL	NUMBER with no precision
RECON_TEST.V_ORG_CHART.DEPTH	NUMBER with no precision

LOW FORCE view R-VIEW-01

PostgreSQL cannot create a view whose dependencies are missing. FORCE views usually signal broken dependencies or circular creation order; both need resolving before the schema replays.

RECON_TEST.V_EMP_DEPT	created with FORCE
RECON_TEST.V_ORG_CHART	created with FORCE

INFO DDL parsed only as an opaque command R-DDL-02

The Oracle grammar accepted the statement without building a real syntax tree, which usually marks exotic constructs the deep-parse pass should revisit.

RECON_TEST.SALES	opaque parse
RECON_TEST.V_EMP_DEPT	opaque parse
RECON_TEST.V_ORG_CHART	opaque parse

INFO Swallowed exception R-SRC-07

WHEN OTHERS THEN NULL hides every failure including the ones a migration introduces. Worth removing before the port so the test phase can actually see errors.

RECON_TEST.ARCHIVE_NOTES	WHEN OTHERS THEN NULL (first at line 22)
--------------------------	--

INFO **Table without optimizer statistics** R-STAT-01

num_rows is null, so data-volume estimates for this table are unreliable. Gather statistics before trusting sizing numbers.

RECON_TEST.AUDIT_LOG	no statistics gathered
RECON_TEST.DEPT	no statistics gathered
RECON_TEST.EMP	no statistics gathered
RECON_TEST.INVOICES	no statistics gathered
RECON_TEST.LEGACY_NOTES	no statistics gathered
RECON_TEST.LEGACY_REFS	no statistics gathered
RECON_TEST.PRODUCT_SPECS	no statistics gathered
RECON_TEST.SALES	no statistics gathered

... and 1 more (9 total; complete list in the findings CSV).

INFO **String lengths declared in bytes** R-TYPE-09

VARCHAR2 and CHAR columns declared without CHAR semantics limit bytes, and in a multibyte database a 30-byte column holds fewer than 30 characters of non-ASCII text. PostgreSQL lengths count characters, so the converted columns accept more text than Oracle did; add CHECK (octet_length(col) <= n) where the byte limit carried meaning, and expect NLS_LENGTH_SEMANTICS to stop mattering.

RECON_TEST.16 string columns	declared in BYTE semantics; PostgreSQL counts characters
------------------------------	--

Risk outlook

The risk profile is driven less by sheer volume and more by the concentration of high-severity constructs that alter transactional and structural behavior, such as autonomous transactions, compound triggers, and package-level state, alongside dependencies like database links and external file access that must be re-architected rather than mechanically translated. The medium-severity population, though larger in count, is generally addressable through established migration patterns but still requires careful validation given the presence of items such as global temporary tables, XMLTYPE columns, and transaction control embedded in PL/SQL. The informational and low findings, including a notable number of tables lacking optimizer statistics, add operational follow-up work but do not by themselves threaten migration feasibility.

Recommended approach

Given the severity mix, with high-impact findings representing a meaningful but minority share of the total against a larger base of medium, low, and informational items, work should be phased to address the high-severity constructs first, since these carry the greatest structural and behavioral risk. Findings are not evenly distributed across the object inventory: a small number of objects account for a disproportionate share of both finding count and weighted complexity, and these should be prioritized early in the program to reduce downstream risk. Subsequent phases can address the medium and low-severity items in a more consolidated fashion, followed by informational cleanup such as statistics gathering. This phased sequencing aligns with the estimated effort range of 106.3 to 180.0 person-days, with 130.9 person-days as the expected case, and supports a controlled, incremental delivery consistent with the program's C grading.

Effort estimate

COMPONENT	PERSON-DAYS
baseline and environment	5.0
schema conversion	1.7
finding remediation	74.3
PL/SQL porting by volume	0.8
data movement	0.0
Development subtotal	81.8

With testing and stabilization: **106 / 131 / 180** person-days (low / expected / high).

Assumptions

- calibration profile 'default': the public engine's conservative defaults, unchanged
- stored code ports at 200 lines per person-day including unit tests
- repeated findings of one rule cost a severity-dependent fraction of the first fix
- testing and stabilization runs 1.3x to 2.2x development, 1.6x expected
- static estimate from schema facts alone: workload, data quality, and application coupling are not visible here

Appendix: findings detail

Each rule lists up to twelve objects here so repeated findings stay readable at estate scale; the findings CSV delivered with this report carries every object with its evidence.

RULE	SEVERITY	OBJECT	EVIDENCE
R-OBJ-01	high	RECON_TEST.LOOPBACK	to //localhost:1521/XEPDB1 as RECON_TEST
R-PART-01	high	RECON_TEST.SALES	INTERVAL NUMTOYMINTERVAL(1, 'MONTH')
R-PERF-02	high	RECON_TEST.SALES_AMOUNT_GIX	GLOBAL partitioned index on SALES
R-PKG-01	high	RECON_TEST.PKG_LEDGER	2 package-level declaration(s), first at line 2
R-PKG-01	high	RECON_TEST.PKG_LEDGER (body)	1 package-level declaration(s), first at line 2
R-SYS-01	high	RECON_TEST.ARCHIVE_NOTES	UTL_FILE (first at line 16)
R-TRG-01	high	RECON_TEST.TRG_EMP_SAL_GUARD	on EMP
R-TRG-02	high	RECON_TEST.TRG_EMP_AUDIT	PRAGMA AUTONOMOUS_TRANSACTION (first at line 5)
R-TYPE-01	high	RECON_TEST.LEGACY_NOTES.BODY	LONG
R-TYPE-07	high	RECON_TEST.LEGACY_REFS.SCAN_DOC	BFILE
R-OBJ-02	medium	RECON_TEST.NIGHTLY_ROLLUP	scheduler job
R-OBJ-03	medium	RECON_TEST.MV_DEPT_SALARIES	materialized view
R-OBJ-05	medium	RECON_TEST.T_MONEY	object type
R-PERF-05	medium	RECON_TEST.MV_DEPT_SALARIES	REWRITE enabled, refresh COMPLETE
R-PKG-02	medium	RECON_TEST.PKG_LEDGER	BEGIN blocks exceed subprograms by 1 (possible initialization block)
R-SRC-01	medium	RECON_TEST.PKG_LEDGER	EXECUTE IMMEDIATE (first at line 27)
R-SRC-02	medium	RECON_TEST.FIND_MANAGER	GOTO (first at line 14)
R-SRC-03	medium	RECON_TEST.PKG_LEDGER	BULK COLLECT (first at line 24)
R-SRC-04	medium	RECON_TEST.V_ORG_CHART	CONNECT BY in view definition
R-SRC-14	medium	RECON_TEST.TRG_EMP_AUDIT	COMMIT or ROLLBACK (first at line 15)
R-SRC-16	medium	RECON_TEST.PKG_LEDGER	collection type declaration (first at line 20)
R-SRC-18	medium	RECON_TEST.ARCHIVE_NOTES	empty-string literal (first at line 4)

R-SRC-19	medium	RECON_TEST.ARCHIVE_NOTES	ROWID (first at line 3)
R-SYS-04	medium	RECON_TEST.ARCHIVE_NOTES	DBMS_LOB (first at line 14)
R-TAB-01	medium	RECON_TEST.STAGING_ROWS	GLOBAL TEMPORARY
R-TYPE-02	medium	RECON_TEST.PRODUCT_SPECS.SPEC	XMLTYPE
R-TYPE-03	medium	RECON_TEST.INVOICES.CREATED_AT	TIMESTAMP(6) WITH LOCAL TIME ZONE
R-TYPE-06	medium	RECON_TEST.LEGACY_REFS.REF_ADDR	ROWID
R-ENV-02	low	RECON_TEST.PUBLIC	1 object grants
R-OBJ-06	low	RECON_TEST.STAFF	for RECON_TEST.EMP
R-PERF-01	low	RECON_TEST.ARCHIVE_NOTES	optimizer hint (first at line 8)
R-PERF-03	low	RECON_TEST.SALES	DEGREE 4
R-SRC-05	low	RECON_TEST.V_EMP_DEPT	uses (+) outer join
R-SRC-11	low	RECON_TEST.ARCHIVE_NOTES	SYSDATE (first at line 19)
R-SRC-11	low	RECON_TEST.PKG_LEDGER	SYSDATE (first at line 7)
R-SRC-11	low	RECON_TEST.TRG_EMP_SAL_GUARD	SYSDATE (first at line 8)
R-SYS-05	low	RECON_TEST.ARCHIVE_NOTES	DBMS_OUTPUT (first at line 19)
R-TYPE-04	low	RECON_TEST.AUDIT_LOG.AUDIT_ID	NUMBER with no precision
R-TYPE-04	low	RECON_TEST.INVOICES.INVOICE_ID	NUMBER with no precision
R-TYPE-04	low	RECON_TEST.MV_DEPT_SALARIES.HEADCOUNT	NUMBER with no precision
R-TYPE-04	low	RECON_TEST.MV_DEPT_SALARIES.TOTAL_SAL	NUMBER with no precision
R-TYPE-04	low	RECON_TEST.V_ORG_CHART.DEPTH	NUMBER with no precision
R-VIEW-01	low	RECON_TEST.V_EMP_DEPT	created with FORCE
R-VIEW-01	low	RECON_TEST.V_ORG_CHART	created with FORCE
R-DDL-02	info	RECON_TEST.SALES	opaque parse
R-DDL-02	info	RECON_TEST.V_EMP_DEPT	opaque parse
R-DDL-02	info	RECON_TEST.V_ORG_CHART	opaque parse
R-SRC-07	info	RECON_TEST.ARCHIVE_NOTES	WHEN OTHERS THEN NULL (first at line 22)
R-STAT-01	info	RECON_TEST.AUDIT_LOG	no statistics gathered
R-STAT-01	info	RECON_TEST.DEPT	no statistics gathered

R-STAT-01	info	RECON_TEST.EMP	no statistics gathered
R-STAT-01	info	RECON_TEST.INVOICES	no statistics gathered
R-STAT-01	info	RECON_TEST.LEGACY_NOTES	no statistics gathered
R-STAT-01	info	RECON_TEST.LEGACY_REFS	no statistics gathered
R-STAT-01	info	RECON_TEST.PRODUCT_SPECS	no statistics gathered
R-STAT-01	info	RECON_TEST.SALES	no statistics gathered
R-STAT-01	info	RECON_TEST.STAGING_ROWS	no statistics gathered
R-TYPE-09	info	RECON_TEST.16 string columns	declared in BYTE semantics; PostgreSQL counts characters

Produced with the pgrecon assessment engine from an offline extraction; no connection was made to the source database. Findings are deterministic and reproducible from the same inventory. DevCrafter · devcrafterai.com